
From Spatial to Spectral: An Efficient, Frequency-Guided Representation Learner for Small Object Detection

Anonymous Authors¹

Abstract

This document provides a basic paper template and submission guidelines. Abstracts must be a single paragraph, ideally between 4–6 sentences long. Gross violations will trigger corrections at the camera-ready phase.

1. Introduction

Small object detection remains a significant challenge due to the limited spatial information they occupy in images. These objects are dominated by high-frequency cues such as edges and fine textures. However, modern detectors face three key issues: (i) Backbone downsampling, which acts as a low-pass filter, reduces these high-frequency components; (ii) Multi-scale fusion in the neck further smooths the boundaries; and (iii) The head’s regression gradients weaken once boundary evidence is diluted during the detection process. These problems make small-object localization particularly difficult.

To address these challenges, we propose a frequency-guided solution that operates on the principle of *decomposing features into low and high-frequency components, selectively enhancing them, and reconstructing or injecting the enhanced signal back into the feature stream*. This approach, encapsulated in the **Frequency-guided Decompose–Enhance–Reconstruct (DER) operator**, enables us to effectively preserve, enhance, and exploit high-frequency details across the detection pipeline.

Building on DER, we introduce three lightweight, plug-and-play components that can be inserted into the backbone, neck, and head, respectively.

These modules form a coherent frequency flow across the backbone, neck, and head, addressing the challenges posed by small-object detection with minimal computational over-

head. Our method is generalizable across different detector architectures, and we demonstrate its effectiveness through extensive experiments across multiple benchmarks.

Contributions: We make the following contributions:

- We propose the DER operator, which integrates frequency-domain decomposition, enhancement, and reconstruction in a unified framework.
- We introduce three novel modules (WDG, LGE, FD-Head) that instantiate the DER operator at key locations in the detection pipeline.
- We show that our method improves small-object detection performance while maintaining computational efficiency, and we validate its generalizability across different architectures.

2. Related Work

We review prior work from three angles that are most relevant to our goal: (i) efficient detector architectures, (ii) small object detection strategies, and (iii) frequency-domain modeling for dense prediction.

2.1. Efficient Detector Architectures

Real-time detection has been driven by architectural efficiency in backbones, feature pyramids, and heads. One-stage YOLO-style detectors optimize the accuracy–latency trade-off through carefully designed blocks and multi-scale prediction, with recent variants continuing to improve both speed and accuracy (Wang et al., 2024; Khanam & Hussain, 2024). Lightweight enhancements for challenging regimes (e.g., cluttered scenes) often rely on stronger feature aggregation or multi-kernel perception to increase representational diversity while keeping inference efficient (Xiao et al., 2025).

In parallel, Transformer-based detectors seek end-to-end set prediction by removing hand-crafted components such as anchors (Carion et al., 2020). Subsequent work improves the practicality of DETR-like models via more efficient attention and training strategies, enabling competitive performance under constrained budgets (Zhao et al., 2024;

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

Zhang et al., 2022). Despite these advances, both CNN- and Transformer-based detectors still face a common tension for tiny/small objects: improving fine-detail sensitivity typically increases computation, memory, or architectural intrusion, making it difficult to deploy a uniformly effective solution across detector families.

2.2. Small Object Detection

Small objects are inherently information-limited: they occupy few pixels, induce weak feature responses, and are easily suppressed by downsampling and coarse fusion. Early two-stage and one-stage frameworks (e.g., Faster R-CNN and SSD) already revealed the difficulty of preserving small-object cues under feature hierarchy and stride growth (Ren et al., 2015; Liu et al., 2016). A large body of work improves small-object performance by strengthening multi-scale feature fusion (e.g., FPN and its variants) (Lin et al., 2017), introducing additional pyramid levels, and designing attention or alignment modules to enhance small-scale features.

Recent methods increasingly emphasize *detail-aware* feature enrichment. For example, HS-FPN highlights tiny objects by generating high-frequency responses as mask weights and complements this with explicit spatial dependency modeling (Shi et al., 2025). Context modeling (e.g., large receptive fields or multi-kernel designs) also helps disambiguate tiny objects from background clutter (Wang et al., 2025; Xiao et al., 2025). However, many of these approaches focus on either spatial fusion or receptive-field engineering, while the *mechanism of how fine details are suppressed and should be reconstructed* is often left implicit, and portability across heterogeneous detector designs is not always validated.

2.3. Frequency-Domain Modeling for Dense Prediction

Frequency-domain analysis offers a complementary lens to understand and manipulate representation learning. A line of work uses Fourier transforms to achieve efficient global interactions. GFNet replaces quadratic self-attention with frequency-domain filtering (FFT-filtering–IFFT), yielding log-linear complexity while maintaining global receptive fields (Rao et al., 2023). Other work links common architectural operations to spectral decomposition: FcaNet interprets channel attention as a frequency-domain compression process and generalizes global pooling to multi-spectral channel attention (?).

More recently, frequency-aware modules have been explored for dense prediction. FDConv observes that candidate dynamic convolution kernels often have highly similar frequency responses, and proposes constructing frequency-diverse weights by allocating parameters to disjoint Fourier indices, together with frequency-band/spatial modulation (Chen et al., 2025). Frequency-aware fusion is also studied:

FreqFusion explicitly introduces adaptive low-pass/high-pass filtering to improve feature consistency and boundary sharpness during upsampling and fusion (?). Wavelet-based approaches provide multi-resolution decomposition with partial spatial localization; WTConv performs convolutions in wavelet sub-bands to scale receptive fields efficiently and can be used as a drop-in layer in CNNs (Finder et al., 2025).

While these spectral methods demonstrate that frequency-domain techniques can be integrated into modern architectures, existing designs are often *task- or component-specific* (e.g., classification backbones, fusion-only modules, or specific convolution families), and do not provide a unified, plug-and-play operator that can be instantiated across *backbone, neck, and head* and generalize across both CNN- and Transformer-style detectors. Our work fills this gap by introducing a decomposition–reconstruction operator that preserves and re-synthesizes discriminative spectral components with minimal overhead, and systematically validating its cross-architecture generality.

3. Method

3.1. Overall Framework

Small object detection is inherently challenging due to the scarcity of spatial information, as small objects occupy few pixels and are dominated by high-frequency cues such as edges and fine textures. However, modern detectors suffer from three key issues: (i) backbone downsampling implicitly acts as a low-pass filter, (ii) multi-scale fusion in the neck smooths boundaries, and (iii) the head’s regression gradients are weakened when boundary information is diluted.

To address these challenges, we propose a frequency-domain solution based on the principle of *decomposing features into low and high-frequency components, selectively enhancing them, and reconstructing or injecting the enhanced signals back into the feature stream*. This solution can be captured by the **Frequency-guided Decompose–Enhance–Reconstruct (DER) operator**.

Frequency-guided Decompose–Enhance–Reconstruct (DER) Operator. Given an input feature tensor $\mathbf{X} \in \mathbb{R}^{C \times H \times W}$, we decompose it into low- and high-frequency components using \mathcal{D} , enhance the components via \mathcal{E}_L and \mathcal{E}_H , and finally reconstruct the enhanced features through \mathcal{R} :

$$\begin{aligned} (\mathbf{X}_L, \mathbf{X}_H) &= \mathcal{D}(\mathbf{X}), \\ \mathbf{X}_L^+ &= \mathcal{E}_L(\mathbf{X}_L), \quad \mathbf{X}_H^+ = \mathcal{E}_H(\mathbf{X}_H), \\ \mathbf{X}^+ &= \mathcal{R}(\mathbf{X}_L^+, \mathbf{X}_H^+). \end{aligned} \quad (1)$$

Here \mathcal{D} extracts low-/high-frequency components (e.g., via wavelet transforms), \mathcal{E}_L and \mathcal{E}_H are lightweight enhancement functions, and \mathcal{R} reconstructs the enhanced signal back into the feature stream.

110 **DER Instantiations Across Backbone, Neck, and Head.**
 111 We instantiate DER at three locations with complementary
 112 roles: WDG in the backbone preserves boundary-relevant
 113 high-frequency evidence before aggressive downsampling;
 114 LGE/LGE-W in the neck re-amplifies high-frequency residuals
 115 before multi-scale fusion to prevent detail dilution; and
 116 FDHead in the head converts high-frequency energy into a
 117 gain factor for boundary-aligned box regression. Together,
 118 they progressively preserve, enhance, and re-inject high-
 119 frequency information with minimal overhead.
 120

121 **Pipeline Overview.** Given a baseline detector with back-
 122 bone \mathcal{B} , neck \mathcal{N} , and head \mathcal{H} , our framework applies the
 123 DER operators as follows:

$$\begin{aligned} \{\mathbf{C}_\ell\} &= \mathcal{B}(\mathbf{I}), \quad \mathbf{C}'_\ell = \begin{cases} \mathcal{W}(\mathbf{C}_\ell), & \ell \in \mathcal{S}_{\mathcal{B}}, \\ \mathbf{C}_\ell, & \text{otherwise}, \end{cases} \\ \{\mathbf{P}_\ell\} &= \mathcal{N}(\{\mathbf{C}'_\ell\}), \quad \mathbf{P}'_\ell = \mathcal{E}(\mathbf{P}_\ell), \\ \widehat{\mathbf{Y}} &= \mathcal{H}_{\text{FD}}(\{\mathbf{P}'_\ell\}). \end{aligned} \quad (2)$$

130 Where \mathcal{W} , \mathcal{E} , and \mathcal{H}_{FD} represent the concrete DER instanti-
 131 ations for the backbone, neck, and head, respectively, and
 132 $\mathcal{S}_{\mathcal{B}}$ denotes the set of backbone stages where WDG is ap-
 133 plied.

3.2. Wavelet-Difference Gate (WDG)

137 We introduce Wavelet-Difference Gate (WDG), a
 138 lightweight plug-and-play bottleneck that injects frequency-
 139 aware modulation into convolutional backbones. Given an
 140 input feature map $\mathbf{x} \in \mathbb{R}^{C \times H \times W}$, WDG first applies a 1×1
 141 projection to hidden channels $C' = \lfloor eC \rfloor$ (with expansion
 142 ratio e) and then performs a 2D Haar discrete wavelet
 143 transform (DWT) to separate low- and high-frequency
 144 components. For simplicity, we describe the transform for
 145 even H, W ; in practice we align sizes by cropping/padding
 146 and restore the original resolution after reconstruction.

147 **Projection and wavelet decomposition.** We first project \mathbf{x}
 148 to a hidden space and decompose it into Haar subbands:
 149

$$\begin{aligned} \mathbf{x}' &= f_{1 \times 1}(\mathbf{x}), \\ (\mathbf{x}_{LL}, \mathbf{x}_{LH}, \mathbf{x}_{HL}, \mathbf{x}_{HH}) &= \text{DWT}(\mathbf{x}'). \end{aligned} \quad (3)$$

153 Here \mathbf{x}_{LL} is the low-frequency approximation, and
 154 $\{\mathbf{x}_{LH}, \mathbf{x}_{HL}, \mathbf{x}_{HH}\}$ capture horizontal/vertical/diagonal
 155 high-frequency details. This decomposition explicitly sepa-
 156 rates coarse structures from fine details, enabling targeted
 157 refinement for small objects.

158 For Haar DWT/IDWT, each spatial 2×2 block is trans-
 159 formed by a 2×2 Haar matrix. For each channel c and
 160 location (u, v) , define the local block
 161

$$\mathbf{X}_{u,v}^{(c)} = \begin{pmatrix} \mathbf{x}_{2u,2v}^{'(c)} & \mathbf{x}_{2u,2v+1}^{'(c)} \\ \mathbf{x}_{2u+1,2v}^{'(c)} & \mathbf{x}_{2u+1,2v+1}^{'(c)} \end{pmatrix}. \quad (4)$$

Then Haar DWT and IDWT are given by

$$\begin{aligned} \mathbf{S}_{u,v}^{(c)} &= \frac{1}{2} \mathbf{H}_2 \mathbf{X}_{u,v}^{(c)} \mathbf{H}_2^\top, \\ \mathbf{X}_{u,v}^{(c)} &= \frac{1}{2} \mathbf{H}_2^\top \mathbf{S}_{u,v}^{(c)} \mathbf{H}_2, \quad \mathbf{H}_2 = \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}, \end{aligned} \quad (5)$$

where $\mathbf{S}_{u,v}^{(c)} = \begin{pmatrix} \mathbf{x}_{LL,u,v}^{(c)} & \mathbf{x}_{LH,u,v}^{(c)} \\ \mathbf{x}_{HL,u,v}^{(c)} & \mathbf{x}_{HH,u,v}^{(c)} \end{pmatrix}$ collects the four sub-
 bands. This matrix form is exactly equivalent to the element-
 wise expressions used in our implementation.

RepCDC for low-frequency refinement. To enhance dis-
 122 criminative edges while keeping computation low, we refine
 the approximation subband at half resolution:

$$\mathbf{y}_{LL} = f_{\text{cdc}}(\mathbf{x}_{LL}). \quad (6)$$

In our implementation, f_{cdc} is RepCDC followed by nor-
 122 malization and activation. RepCDC parameterizes a central-
 difference convolution by decreasing the center coefficient
 of a 3×3 kernel with a learnable $\boldsymbol{\theta}$. Concretely, the effective
 kernel is

$$\mathbf{y}_{p,q}^{(o)} = \sum_c \sum_{i=-1}^1 \sum_{j=-1}^1 \mathbf{W}_{i,j}^{(o,c)} \mathbf{z}_{p+i,q+j}^{(c)} - \sum_c \boldsymbol{\theta}^{(o,c)} \mathbf{z}_{p,q}^{(c)}, \quad (7)$$

where \mathbf{z} denotes the input to RepCDC (e.g., $\mathbf{z} = \mathbf{x}_{LL}$), and
 122 (p, q) indexes spatial locations. This expression is exactly
 equivalent to subtracting $\boldsymbol{\theta}$ from the center coefficient of a
 123 3×3 kernel. During deployment, the resulting kernel is
 fused into a single standard convolution, so RepCDC incurs
 no extra inference branches. Operating on \mathbf{x}_{LL} reduces spa-
 tial cost by $4 \times$ while strengthening edge sensitivity through
 the difference term.

High-frequency gated modulation. We use high-frequency
 122 responses to predict a content-adaptive gate and modulate
 the refined low-frequency feature:

$$\begin{aligned} \mathbf{g} &= \sigma(f_g(\text{Concat}(\mathbf{x}_{LH}, \mathbf{x}_{HL}, \mathbf{x}_{HH}))), \\ \widetilde{\mathbf{x}}_{LL} &= \mathbf{y}_{LL} \odot (\mathbf{1} + \mathbf{g}). \end{aligned} \quad (8)$$

We use additive gating $(\mathbf{1} + \mathbf{g})$ to emphasize informative
 122 regions without suppressing the overall magnitude of \mathbf{y}_{LL} .
 f_g is a 1×1 convolution followed by normalization, and
 123 $\text{Concat}(\cdot)$ denotes channel-wise concatenation. Since the
 gate is predicted from high-frequency subbands, it acts
 as a detail-aware selector that boosts regions with strong
 edge/texture cues.

Reconstruction and residual output. Finally, we preserve
 the original high-frequency subbands and reconstruct the
 feature via inverse Haar transform:

$$\begin{aligned} \widehat{\mathbf{x}}' &= \text{IDWT}(\widetilde{\mathbf{x}}_{LL}, \mathbf{x}_{LH}, \mathbf{x}_{HL}, \mathbf{x}_{HH}), \\ \mathbf{y} &= f_{1 \times 1}^{\text{out}}(\widehat{\mathbf{x}}'). \end{aligned} \quad (9)$$

When input/output channels match, WDG uses a residual connection $\mathbf{y} \leftarrow \mathbf{x} + \mathbf{y}$. Since the wavelet-domain refinement operates on $H/2 \times W/2$, WDG adds only a small overhead and can be inserted as a generic bottleneck into different backbone designs. Preserving the original high-frequency subbands avoids over-smoothing and helps retain boundary sharpness after reconstruction.

3.3. Log-Gabor Enhancer (LGE) and WTConv Variant (LGE-W)

We next improve the neck by introducing Log-Gabor Enhancer (LGE), a plug-and-play high-frequency refinement module applied to intermediate feature maps before multi-scale fusion. LGE is instantiated per feature level and is agnostic to the specific fusion topology (e.g., FPN/PAN/decoder-style aggregation).

Log-Gabor filter bank (LGF). Given a feature map $\mathbf{x} \in \mathbb{R}^{C \times H \times W}$, LGF applies a fixed Log-Gabor filter bank using depthwise convolutions. Let K and S denote the number of orientations and scales. For each channel c , orientation k , and scale s , we compute

$$\mathbf{h}_{s,k}^{(c)} = \mathbf{x}^{(c)} * \mathbf{g}_{s,k}, \quad (10)$$

where $\mathbf{g}_{s,k}$ is a non-learnable Log-Gabor kernel and $*$ is convolution. In our implementation, $\mathbf{g}_{s,k}$ is instantiated in the spatial domain by rotating a centered coordinate grid and applying a log-normal radial envelope with a cosine angular term:@@

$$\begin{aligned} c_k &= \cos \phi_k, & s_k &= \sin \phi_k, \\ u' &= u c_k + v s_k, & v' &= -u s_k + v c_k, \\ r &= \sqrt{u'^2 + v'^2} + \varepsilon, & \theta &= \text{atan2}(v', u'), \\ \mathbf{g}_{s,k}(u, v) &= \exp\left(-\frac{\log^2(r/\rho_s)}{2 \log^2 2}\right) \cos \theta. \end{aligned} \quad (11)$$

where $\phi_k = k\pi/K$ and ρ_s is a fixed scale parameter. This produces a set of directional subband responses that explicitly emphasize edges and fine textures while introducing no additional learnable filter parameters.

Learnable aggregation and residual enhancement (LGE). LGE aggregates the subbands with learnable orientation/scale importance. Let $\boldsymbol{\alpha} \in \mathbb{R}^S$ and $\boldsymbol{\beta} \in \mathbb{R}^K$ be learnable logits; we obtain normalized weights by softmax and compute the high-frequency summary

$$\mathbf{h}^{(c)} = \sum_{s=1}^S \sum_{k=1}^K \text{softmax}(\boldsymbol{\alpha})_s \text{softmax}(\boldsymbol{\beta})_k \mathbf{h}_{s,k}^{(c)}. \quad (12)$$

We further apply a learnable global scale γ (implemented as a sigmoid-gated parameter) and a local mixing operator f_{mix} :

$$\mathbf{y} = \mathbf{x}_{\text{skip}} + f_{\text{mix}}(\sigma(\gamma) \mathbf{h}). \quad (13)$$

Here \mathbf{x}_{skip} is either the identity mapping (when channels match) or a 1×1 projection. In our implementation, f_{mix} is a 3×3 convolution (depthwise when C is preserved), so LGE adds only local mixing on top of fixed spectral decomposition while keeping a residual pathway.

Wavelet variant (LGE-W). LGE-W follows Eq. (10)–(13) but replaces f_{mix} with a wavelet-transform convolution (WTConv) when C is preserved. Using a fixed wavelet (Haar/db1), WTConv performs subband mixing in the wavelet domain and adds a lightweight depthwise branch:@@

$$\text{WTConv}(\mathbf{z}) = \mathcal{S}_0 \mathcal{D}_0(\mathbf{z}) + \text{IDWT}(\mathcal{S} \mathcal{D}_4(\text{DWT}(\mathbf{z}))), \quad (14)$$

where \mathcal{D}_0 is a depthwise convolution in the spatial domain and \mathcal{D}_4 denotes grouped depthwise convolutions applied over the four wavelet subbands. @@

3.4. Frequency-Driven Head (FDHead)

We finally introduce Frequency-Driven Head (FDHead), a frequency-aware detection head that improves small-object localization by injecting a boundary-sensitive prior into dense regression while preserving the standard anchor-free interface. FDHead is instantiated over multi-scale feature maps $\{\mathbf{x}_i\}_{i=1}^N$ (e.g., $P2-P5$) and shares most head parameters across levels to reduce capacity fragmentation.

Shared prediction tower. For each level i , FDHead first aligns channels to a hidden width C_h (Conv+GroupNorm) and then applies a shared refinement stack (DEConv + depthwise-pointwise mixing). The DEConv block aggregates multiple directional-difference operators (center/adjacent/horizontal/vertical) and a standard kernel; at inference it can be written as a single convolution with merged weights:

$$\text{DEConv}(\mathbf{u}) = \varphi\left(\left(\sum_m \mathbf{K}_m\right) * \mathbf{u} + \sum_m \mathbf{b}_m\right), \quad (15)$$

where m indexes the directional branches and $\varphi(\cdot)$ denotes normalization and activation. This biases the shared tower toward contour-aware features that are beneficial for boundary-aligned regression.

$$\mathbf{f}_i = \mathcal{T}(\mathbf{x}_i), \quad \mathcal{T} = \mathcal{T}_{\text{share}} \circ \mathcal{T}_{1 \times 1}. \quad (16)$$

P2 high-frequency gate. Since the finest level ($P2$) carries the most precise spatial details, FDHead applies a lightweight wavelet gate only on $i = 1$ (corresponding to $P2$). Let C_f be the gated channel width (set as a fraction of C_h); we split channels $\mathbf{f}_1 = [\mathbf{f}_a, \mathbf{f}_b]$ with $\mathbf{f}_a \in \mathbb{R}^{C_f \times H \times W}$. Using a fixed Haar transform, we estimate high-frequency energy as a softmax-weighted mixture of subband magni-

tudes and convert it to a channel-wise gain:

$$\begin{aligned}
 (\mathbf{f}_{LL}, \mathbf{f}_{LH}, \mathbf{f}_{HL}, \mathbf{f}_{HH}) &= \text{DWT}(\mathbf{f}_a), \\
 \mathbf{w} &= \text{softmax}(\boldsymbol{\omega}), \\
 \mathbf{h} &= w_{LH} |\mathbf{f}_{LH}| + w_{HL} |\mathbf{f}_{HL}| + w_{HH} |\mathbf{f}_{HH}|, \\
 \mathbf{g} &= \text{Gate}(\text{AvgPool}(\mathbf{h})), \\
 \tilde{\mathbf{f}}_a &= \mathbf{f}_a \odot (1 + \alpha \mathbf{g}). \tag{17}
 \end{aligned}$$

Here $\boldsymbol{\omega}$ are learnable logits over $\{LH, HL, HH\}$ and α controls the gate strength. Gate(\cdot) is a squeeze-excitation style channel MLP (two 1×1 convs with sigmoid output) driven by pooled high-frequency energy. We then form $\tilde{\mathbf{f}}_1 = [\tilde{\mathbf{f}}_a, \mathbf{f}_b]$ and apply it only to the box branch: high-frequency energy is a direct proxy for boundary sharpness and thus improves offset estimation, while leaving the classification stream unchanged avoids over-fitting to textures and background clutter. For the remaining levels $i > 1$, we set $\tilde{\mathbf{f}}_i = \mathbf{f}_i$.

Box/class prediction and decoding. FDHead predicts per-location class logits and distributional box offsets (DFL) as

$$\mathbf{b}_i = \text{Scale}_i(\mathcal{H}_{\text{box}}(\tilde{\mathbf{f}}_i)), \quad \mathbf{p}_i = \mathcal{H}_{\text{cls}}(\mathbf{f}_i), \tag{18}$$

and decodes boxes by $\hat{\mathbf{B}} = \text{dist2bbox}(\text{DFL}(\mathbf{b}), \mathbf{A}) \cdot \mathbf{s}$ with anchors \mathbf{A} and strides \mathbf{s} . This design targets small objects by frequency-gating only the finest level while keeping the remaining head computation shared and lightweight.

4. Experiment

4.1. Datasets and Metrics

We evaluate our framework on four benchmarks to demonstrate its robustness and cross-domain generalization: VisDrone2019 (Du et al., 2019), TinyPerson (Yu et al., 2020), UAVDT (Du et al., 2018), and DOTA v1 (Xia et al., 2018). **VisDrone2019** is our primary benchmark and is particularly challenging due to dense small objects and severe scale variation, where most targets are smaller than 50×50 pixels.

We report both accuracy and efficiency, including mAP₅₀, the number of parameters, GFLOPs, model size, and FPS.

4.2. Configuration

The experimental configuration is detailed in Table 1.

For YOLO-style architectures, models are trained for 300 epochs with an input resolution of 640×640 and batch size 16, using SGD optimization. Unless otherwise specified, Mosaic augmentation is enabled throughout training; we use 4 dataloader workers and disable AMP.

Table 1. Configuration of Training and Testing Experiment Environments. Detailed hardware and software configuration used for all experiments in this study.

Environment	Parameter
CPU	Intel(R) Xeon(R) Gold 5218R CPU @ 2.10GHz
GPU	NVIDIA A100-PCIE-40GB
VRAM	40 GB
RAM	46 GB
Operating System	Rocky Linux 8.5 (Green Obsidian)
Language	Python 3.10.14
Frame	PyTorch 2.1.0
CUDA Version	12.6

Top three performances per metric are highlighted from Dark (best) to Light (third).

P2 head	WDG	LGE	FDHead	Performance						Efficiency			
				Val			Test			Params/M	GFLOPs	FPS	Model size/MB
✗	✗	✗	✗	0.496	0.377	0.384	0.421	0.337	0.311	9.4	21.6	–	18.3
✓	✗	✗	✗	0.516	0.408	0.423	0.453	0.354	0.337	6.4	24.7	–	12.5
✓	✓	✗	✗	0.543	0.420	0.438	0.455	0.371	0.354	6.6	23.5	–	13.2
✓	✗	✓	✗	0.547	0.412	0.436	0.465	0.373	0.355	6.5	25.4	–	12.8
✓	✓	✗	✓	0.551	0.435	0.454	0.478	0.375	0.359	3.8	30.5	–	8.4
✓	✓	✓	✗	0.547	0.414	0.436	0.469	0.370	0.360	3.9	25.5	–	8.2
✓	✓	✓	✓	0.559	0.435	0.464	0.487	0.383	0.370	3.9	28.5	–	8.6
✓	✓	✓	✓	0.549	0.437	0.457	0.470	0.378	0.365	3.8	28.3	–	8.4
✓	✓	✓	✓	0.552	0.429	0.458	0.484	0.381	0.370	3.8	26.3	–	8.6
✓	✓	Lite	✓	0.503	0.382	0.398	0.420	0.344	0.316	1.3	13.3	–	3.4

Table 2. Ablation study on YOLO-style architectures.

5. Main Results

5.1. Ablation Study on YOLO-style architectures

5.2. Across-architecture Study

5.3. Comparison with State-of-the-art

6. Analyses and Discussion

7. Conclusion

References

- Carion, N., Massa, F., Synnaeve, G., Usunier, N., Kirillov, A., and Zagoruyko, S. End-to-end object detection with transformers. *arXiv preprint arXiv:2005.12872*, 2020.
- Chen, L., Gu, L., Li, L., Yan, C., and Fu, Y. Frequency dynamic convolution for dense image prediction. *arXiv preprint arXiv:2503.18783*, 2025.
- Du, D., Qi, Y., Yu, H., Yang, Y., Duan, K., Li, G., Zhang, W., Huang, Q., and Tian, Q. The unmanned aerial vehicle benchmark: Object detection and tracking. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pp. 370–386, 2018.
- Du, D., Zhu, P., Wen, L., Bian, X., Lin, H., Hu, Q., Peng, T., Zheng, J., Wang, X., Zhang, Y., et al. Visdrone-det2019: The vision meets drone object detection in image challenge results. In *Proceedings of the IEEE/CVF International Conference on Computer Vision Workshops*, pp. 0–0, 2019.
- Finder, S. E., Amoyal, R., Treister, E., and Freifeld, O. Wavelet convolutions for large receptive fields. In

- 275 Leonardis, A., Ricci, E., Roth, S., Russakovsky, O., Sat-
 276 tter, T., and Varol, G. (eds.), *Computer Vision – ECCV*
 277 2024, pp. 363–380, Cham, 2025. Springer Nature Switzer-
 278 land. ISBN 978-3-031-72949-2.
- 279
 280 Khanam, R. and Hussain, M. YOLOv11: An overview
 281 of the key architectural enhancements. *arXiv preprint*
 282 *arXiv:2410.17725*, 2024.
- 283
 284 Lin, T.-Y., Dollár, P., Girshick, R., He, K., Hariharan, B.,
 285 and Belongie, S. Feature pyramid networks for object
 286 detection. In *Proceedings of the IEEE Conference on*
 287 *Computer Vision and Pattern Recognition (CVPR)*, pp.
 288 2117–2125, 2017.
- 289
 290 Liu, W., Anguelov, D., Erhan, D., Szegedy, C., Reed, S.,
 291 Fu, C.-Y., and Berg, A. C. SSD: Single shot multibox
 292 detector. In *Computer Vision–ECCV 2016: 14th Euro-
 293 pean Conference, Amsterdam, The Netherlands, October
 294 11–14, 2016, Proceedings, Part I*, pp. 21–37. Springer,
 295 2016.
- 296
 297 Rao, Y., Zhao, W., Zhu, Z., Zhou, J., and Lu, J. GFNet:
 298 Global filter networks for visual recognition. *IEEE Trans-
 299 actions on Pattern Analysis and Machine Intelligence
 300 (TPAMI)*, 45(9):10960–10973, September 2023. doi:
 301 10.1109/TPAMI.2023.3263824.
- 302
 303 Ren, S., He, K., Girshick, R., and Sun, J. Faster R-CNN:
 304 Towards real-time object detection with region proposal
 305 networks. In *Advances in Neural Information Processing
 306 Systems (NeurIPS)*, volume 28, 2015.
- 307
 308 Shi, Z., Hu, J., Ren, J., Ye, H., Yuan, X., Ouyang, Y., He, J.,
 309 Ji, B., and Guo, J. HS-FPN: High frequency and spatial
 310 perception fpn for tiny object detection. *arXiv preprint*
arXiv:2412.10116, 2025.
- 311
 312 Wang, A., Chen, H., Liu, L., Chen, K., Lin, Z., Han, J.,
 313 and Ding, G. YOLOv10: Real-time end-to-end object
 314 detection. *arXiv preprint arXiv:2405.14458*, 2024.
- 315
 316 Wang, A., Chen, H., Lin, Z., Han, J., and Ding, G. LSNet:
 317 See large, focus small. In *Proceedings of the IEEE/CVF*
 318 *Conference on Computer Vision and Pattern Recognition
 319 (CVPR)*, 2025.
- 320
 321 Xia, G.-S., Bai, X., Ding, J., Zhu, Z., Belongie, S., Luo,
 322 J., Datcu, M., Pelillo, M., and Zhang, L. Dota: A large-
 323 scale dataset for object detection in aerial images. In
 324 *Proceedings of the IEEE conference on computer vision
 325 and pattern recognition*, pp. 3974–3983, 2018.
- 326
 327 Xiao, Y., Xu, T., Xin, Y., and Li, J. FBRT-YOLO: Faster and
 328 better for real-time aerial image detection. *arXiv preprint*
arXiv:2504.20670, 2025.
- 329
 Yu, X., Gong, Y., Jiang, N., Ye, Q., and Han, Z. Scale
 match for tiny person detection. In *Proceedings of the
 IEEE/CVF Winter Conference on Applications of Com-
 puter Vision*, pp. 1257–1265, 2020.
- Zhang, H., Li, F., Liu, S., Zhang, L., Su, H., Zhu, J., Ni,
 L. M., and Shum, H.-Y. DINO: DETR with improved
 denoising anchor boxes for end-to-end object detection.
arXiv preprint arXiv:2203.03605, 2022.
- Zhao, Y., Lv, W., Xu, S., Wei, J., Wang, G., Dang, Q., Liu,
 Y., and Chen, J. Detrs beat yolos on real-time object
 detection. *arXiv preprint arXiv:2304.08069*, 2024.

330 **A. You *can* have an appendix here.**

331 You can have as much text here as you want. The main body must be at most 8 pages long. For the final version, one more
332 page can be added. If you want, you can use an appendix like this one.
333

334 The `\onecolumn` command above can be kept in place if you prefer a one-column appendix, or can be removed if you
335 prefer a two-column appendix. Apart from this possible change, the style (font size, spacing, margins, page numbering, etc.)
336 should be kept the same as the main body.
337

338

339

340

341

342

343

344

345

346

347

348

349

350

351

352

353

354

355

356

357

358

359

360

361

362

363

364

365

366

367

368

369

370

371

372

373

374

375

376

377

378

379

380

381

382

383

384